

Homework 4: Build a CNN for Image Recognition

Name: Tyler Bryk

1. Data Preparation

1.1. Load data

```
In [1]: from keras.datasets import cifar10
import numpy

(x_train, y_train), (x_test, y_test) = cifar10.load_data()

print('shape of x_train: ' + str(x_train.shape))
print('shape of y_train: ' + str(y_train.shape))
print('shape of x_test: ' + str(x_test.shape))
print('shape of y_test: ' + str(y_test.shape))
print('number of classes: ' + str(numpy.max(y_train) - numpy.min(y_train) + 1))
```

Using TensorFlow backend.

```
shape of x_train: (50000, 32, 32, 3)
shape of y_train: (50000, 1)
shape of x_test: (10000, 32, 32, 3)
shape of y_test: (10000, 1)
number of classes: 10
```

1.2. One-hot encode the labels

In the input, a label is a scalar in $\{0, 1, \dots, 9\}$. One-hot encode transform such a scalar to a 10-dim vector. E.g., a scalar `y_train[j]=3` is transformed to the vector `y_train_vec[j]= $[0, 0, 0, 1, 0, 0, 0, 0, 0, 0]$` .

- Define a function `to_one_hot` that transforms an $n \times 1$ array to a $n \times 10$ matrix.
- Apply the function to `y_train` and `y_test`.

```
In [2]: def to_one_hot(y, num_class=10):
        results = numpy.zeros((len(y), num_class))
        for i, label in enumerate(y):
            results[i, label] = 1.
        return results

y_train_vec = to_one_hot(y_train)
y_test_vec = to_one_hot(y_test)

print('Shape of y_train_vec: ' + str(y_train_vec.shape))
print('Shape of y_test_vec: ' + str(y_test_vec.shape))

print(y_train[0])
print(y_train_vec[0])

Shape of y_train_vec: (50000, 10)
Shape of y_test_vec: (10000, 10)
[6]
[0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]

In [3]: # Convert Image Data to Float-32
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

# Standardize Data with Z-scores
mean = numpy.mean(x_train, axis=(0,1,2,3))
std = numpy.std(x_train, axis=(0,1,2,3))
x_train = (x_train - mean) / (std + 1e-7)
x_test = (x_test - mean) / (std + 1e-7)
```

1.3. Randomly partition the training set to training and validation sets

Randomly partition the 50K training samples to 2 sets:

- a training set containing 40K samples
- a validation set containing 10K samples

```
In [4]: rand_indices = numpy.random.permutation(50000)
train_indices = rand_indices[0:40000]
valid_indices = rand_indices[40000:50000]

x_val = x_train[valid_indices, :]
y_val = y_train_vec[valid_indices, :]

x_tr = x_train[train_indices, :]
y_tr = y_train_vec[train_indices, :]
```

```
print('Shape of x_tr: ' + str(x_tr.shape))
print('Shape of y_tr: ' + str(y_tr.shape))
print('Shape of x_val: ' + str(x_val.shape))
print('Shape of y_val: ' + str(y_val.shape))

Shape of x_tr: (40000, 32, 32, 3)
Shape of y_tr: (40000, 10)
Shape of x_val: (10000, 32, 32, 3)
Shape of y_val: (10000, 10)
```

2. Build a CNN and tune its hyper-parameters

- Build a convolutional neural network model
- Use the validation data to tune the hyper-parameters (e.g., network structure, and optimization algorithm)
 - Do NOT use test data for hyper-parameter tuning!!!
- Try to achieve a validation accuracy as high as possible.

```
In [5]: from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, Activation
from keras.layers.normalization import BatchNormalization
from keras.models import Sequential
from keras.regularizers import l2

model = Sequential()
model.add(Conv2D(32, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu', input_shape=(32,32,3)))
model.add(BatchNormalization())
model.add(Conv2D(32, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(64, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Dropout(0.3))
model.add(Conv2D(128, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu'))
model.add(BatchNormalization())
model.add(Conv2D(128, (3, 3), padding='same', kernel_regularizer=l2(1e-4), activation='relu'))
model.add(BatchNormalization())
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Flatten())
model.add(Dense(512))
model.add(BatchNormalization())
model.add(Activation('relu'))
model.add(Dropout(0.4))
model.add(Dense(10))
model.add(BatchNormalization())
model.add(Activation('softmax'))

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_2 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_2 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_1 (Dropout)	(None, 16, 16, 32)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_4 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_2 (Dropout)	(None, 8, 8, 64)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_5 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_6 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_6 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten_1 (Flatten)	(None, 2048)	0
dense_1 (Dense)	(None, 512)	1049088
batch_normalization_7 (Batch Normalization)	(None, 512)	2048
activation_1 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130
batch_normalization_8 (Batch Normalization)	(None, 10)	40
activation_2 (Activation)	(None, 10)	0

Total params: 1,345,106
Trainable params: 1,343,166
Non-trainable params: 1,940

```
In [6]: from keras import optimizers
from keras.preprocessing.image import ImageDataGenerator

gen = ImageDataGenerator(rotation_range=15,
                          width_shift_range=0.1,
                          height_shift_range=0.1,
                          shear_range=0.3,
                          zoom_range=0.1,
                          horizontal_flip=True)

gen.fit(x_tr)
train_generator = gen.flow(x_tr, y_tr, batch_size=64)
test_generator = ImageDataGenerator().flow(x_val, y_val, batch_size=64)

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.RMSprop(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

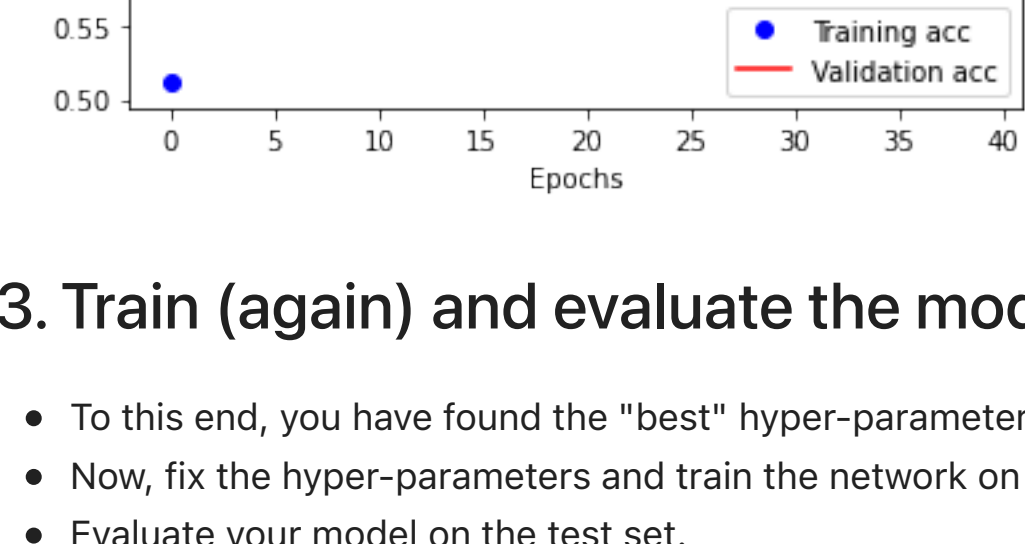
history = model.fit_generator(train_generator, steps_per_epoch=40000//64, epochs=40, validation_data=test_generator, validation_steps=10000//64)

Epoch 1/40
625/625 [=====] - 43s 69ms/step - loss: 1.4541 - accuracy: 0.5121 - val_loss: 0.9724 - val_accuracy: 0.6383
Epoch 2/40
625/625 [=====] - 42s 67ms/step - loss: 1.1017 - accuracy: 0.6489 - val_loss: 0.8180 - val_accuracy: 0.7078
Epoch 3/40
625/625 [=====] - 43s 69ms/step - loss: 0.9649 - accuracy: 0.6999 - val_loss: 1.0764 - val_accuracy: 0.7241
Epoch 4/40
625/625 [=====] - 41s 66ms/step - loss: 0.8882 - accuracy: 0.7270 - val_loss: 0.7428 - val_accuracy: 0.7729
Epoch 5/40
625/625 [=====] - 42s 67ms/step - loss: 0.8387 - accuracy: 0.7444 - val_loss: 0.8329 - val_accuracy: 0.7653
Epoch 6/40
625/625 [=====] - 42s 67ms/step - loss: 0.8049 - accuracy: 0.7580 - val_loss: 0.8565 - val_accuracy: 0.7665
Epoch 7/40
625/625 [=====] - 42s 67ms/step - loss: 0.7808 - accuracy: 0.7675 - val_loss: 0.7876 - val_accuracy: 0.7876
Epoch 8/40
625/625 [=====] - 42s 67ms/step - loss: 0.7617 - accuracy: 0.7760 - val_loss: 0.7773 - val_accuracy: 0.8074
Epoch 9/40
625/625 [=====] - 42s 66ms/step - loss: 0.7474 - accuracy: 0.7815 - val_loss: 0.8030 - val_accuracy: 0.8002
Epoch 10/40
625/625 [=====] - 42s 66ms/step - loss: 0.7249 - accuracy: 0.7909 - val_loss: 0.5733 - val_accuracy: 0.8044
Epoch 11/40
625/625 [=====] - 42s 66ms/step - loss: 0.7145 - accuracy: 0.7949 - val_loss: 0.4619 - val_accuracy: 0.8152
Epoch 12/40
625/625 [=====] - 42s 67ms/step - loss: 0.7036 - accuracy: 0.7997 - val_loss: 0.4383 - val_accuracy: 0.8076
Epoch 13/40
625/625 [=====] - 42s 67ms/step - loss: 0.6994 - accuracy: 0.8030 - val_loss: 0.5725 - val_accuracy: 0.8365
Epoch 14/40
625/625 [=====] - 42s 66ms/step - loss: 0.6889 - accuracy: 0.8055 - val_loss: 0.4632 - val_accuracy: 0.8008
Epoch 15/40
625/625 [=====] - 41s 66ms/step - loss: 0.6746 - accuracy: 0.8124 - val_loss: 0.5892 - val_accuracy: 0.8456
Epoch 16/40
625/625 [=====] - 42s 67ms/step - loss: 0.6696 - accuracy: 0.8156 - val_loss: 0.5663 - val_accuracy: 0.8384
Epoch 17/40
625/625 [=====] - 41s 66ms/step - loss: 0.6704 - accuracy: 0.8146 - val_loss: 0.4523 - val_accuracy: 0.8565
Epoch 18/40
625/625 [=====] - 41s 66ms/step - loss: 0.6605 - accuracy: 0.8184 - val_loss: 0.6145 - val_accuracy: 0.8458
Epoch 19/40
625/625 [=====] - 41s 66ms/step - loss: 0.6567 - accuracy: 0.8209 - val_loss: 0.6421 - val_accuracy: 0.8353
Epoch 20/40
625/625 [=====] - 41s 66ms/step - loss: 0.6596 - accuracy: 0.8223 - val_loss: 0.3942 - val_accuracy: 0.8530
Epoch 21/40
625/625 [=====] - 41s 66ms/step - loss: 0.6466 - accuracy: 0.8251 - val_loss: 0.6807 - val_accuracy: 0.8385
Epoch 22/40
625/625 [=====] - 42s 66ms/step - loss: 0.6471 - accuracy: 0.8229 - val_loss: 0.5245 - val_accuracy: 0.8259
Epoch 23/40
625/625 [=====] - 41s 66ms/step - loss: 0.6419 - accuracy: 0.8288 - val_loss: 0.8233 - val_accuracy: 0.8298
Epoch 24/40
625/625 [=====] - 41s 66ms/step - loss: 0.6336 - accuracy: 0.8307 - val_loss: 0.4700 - val_accuracy: 0.8462
Epoch 25/40
625/625 [=====] - 41s 66ms/step - loss: 0.6388 - accuracy: 0.8276 - val_loss: 0.5480 - val_accuracy: 0.8606
Epoch 26/40
625/625 [=====] - 42s 67ms/step - loss: 0.6337 - accuracy: 0.8329 - val_loss: 0.6637 - val_accuracy: 0.8364
Epoch 27/40
625/625 [=====] - 42s 67ms/step - loss: 0.6286 - accuracy: 0.8335 - val_loss: 0.5554 - val_accuracy: 0.8480
Epoch 28/40
625/625 [=====] - 42s 66ms/step - loss: 0.6285 - accuracy: 0.8339 - val_loss: 0.5141 - val_accuracy: 0.8493
Epoch 29/40
625/625 [=====] - 41s 66ms/step - loss: 0.6237 - accuracy: 0.8366 - val_loss: 0.8269 - val_accuracy: 0.8449
Epoch 30/40
625/625 [=====] - 41s 66ms/step - loss: 0.6183 - accuracy: 0.8379 - val_loss: 0.6516 - val_accuracy: 0.8405
Epoch 31/40
625/625 [=====] - 41s 66ms/step - loss: 0.6176 - accuracy: 0.8383 - val_loss: 0.5224 - val_accuracy: 0.8589
Epoch 32/40
625/625 [=====] - 42s 67ms/step - loss: 0.6138 - accuracy: 0.8399 - val_loss: 0.9207 - val_accuracy: 0.8479
Epoch 33/40
625/625 [=====] - 42s 67ms/step - loss: 0.6120 - accuracy: 0.8411 - val_loss: 0.4478 - val_accuracy: 0.8698
Epoch 34/40
625/625 [=====] - 42s 67ms/step - loss: 0.6134 - accuracy: 0.8412 - val_loss: 0.5255 - val_accuracy: 0.8681
Epoch 35/40
625/625 [=====] - 42s 67ms/step - loss: 0.6062 - accuracy: 0.8421 - val_loss: 0.6383 - val_accuracy: 0.8624
Epoch 36/40
625/625 [=====] - 42s 67ms/step - loss: 0.6113 - accuracy: 0.8414 - val_loss: 0.4884 - val_accuracy: 0.8503
Epoch 37/40
625/625 [=====] - 41s 66ms/step - loss: 0.6023 - accuracy: 0.8436 - val_loss: 0.5688 - val_accuracy: 0.8629
Epoch 38/40
625/625 [=====] - 42s 67ms/step - loss: 0.5996 - accuracy: 0.8451 - val_loss: 0.5783 - val_accuracy: 0.8639
Epoch 39/40
625/625 [=====] - 41s 66ms/step - loss: 0.5999 - accuracy: 0.8450 - val_loss: 0.3326 - val_accuracy: 0.8593
Epoch 40/40
625/625 [=====] - 41s 66ms/step - loss: 0.5995 - accuracy: 0.8454 - val_loss: 0.5971 - val_accuracy: 0.8702
```

```
In [7]: import matplotlib.pyplot as plt
```

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(len(acc))

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'r', label='Validation acc')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



3. Train (again) and evaluate the model

- To this end, you have found the "best" hyper-parameters.
- Now, fix the hyper-parameters and train the network on the entire training set (all the 50K training samples)
- Evaluate your model on the test set.

3.1. Train the model on the entire training set

```
In [8]: gen = ImageDataGenerator(rotation_range=15,
                              width_shift_range=0.1,
                              height_shift_range=0.1,
                              shear_range=0.3,
                              zoom_range=0.1,
                              horizontal_flip=True)

gen.fit(x_train)
train_generator = gen.flow(x_train, y_train_vec, batch_size=64)

model.compile(loss='categorical_crossentropy',
              optimizer=optimizers.RMSprop(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

history = model.fit_generator(train_generator, steps_per_epoch=50000//64, epochs=40)

Epoch 1/40
781/781 [=====] - 52s 66ms/step - loss: 0.6239 - accuracy: 0.8400
Epoch 2/40
781/781 [=====] - 50s 64ms/step - loss: 0.6187 - accuracy: 0.8405
Epoch 3/40
781/781 [=====] - 50s 64ms/step - loss: 0.6162 - accuracy: 0.8402
Epoch 4/40
781/781 [=====] - 50s 65ms/step - loss: 0.6115 - accuracy: 0.8442
Epoch 5/40
781/781 [=====] - 51s 65ms/step - loss: 0.6103 - accuracy: 0.8436
Epoch 6/40
781/781 [=====] - 51s 66ms/step - loss: 0.6114 - accuracy: 0.8461
Epoch 7/40
781/781 [=====] - 49s 63ms/step - loss: 0.6091 - accuracy: 0.8450
Epoch 8/40
781/781 [=====] - 50s 64ms/step - loss: 0.6091 - accuracy: 0.8459
Epoch 9/40
781/781 [=====] - 50s 65ms/step - loss: 0.5970 - accuracy: 0.8490
Epoch 10/40
781/781 [=====] - 50s 64ms/step - loss: 0.6001 - accuracy: 0.8487
Epoch 11/40
781/781 [=====] - 50s 64ms/step - loss: 0.5925 - accuracy: 0.8507
Epoch 12/40
781/781 [=====] - 50s 64ms/step - loss: 0.5953 - accuracy: 0.8490
Epoch 13/40
781/781 [=====] - 50s 64ms/step - loss: 0.6024 - accuracy: 0.8488
Epoch 14/40
781/781 [=====] - 50s 64ms/step - loss: 0.5903 - accuracy: 0.8526
Epoch 15/40
781/781 [=====] - 50s 65ms/step - loss: 0.5881 - accuracy: 0.8502
Epoch 16/40
781/781 [=====] - 50s 65ms/step - loss: 0.5758 - accuracy: 0.8572
Epoch 17/40
781/781 [=====] - 50s 64ms/step - loss: 0.5884 - accuracy: 0.8510
Epoch 18/40
781/781 [=====] - 50s 64ms/step - loss: 0.5894 - accuracy: 0.8512
Epoch 19/40
781/781 [=====] - 50s 64ms/step - loss: 0.5878 - accuracy: 0.8525
Epoch 20/40
781/781 [=====] - 50s 64ms/step - loss: 0.5820 - accuracy: 0.8562
Epoch 21/40
781/781 [=====] - 50s 64ms/step - loss: 0.5816 - accuracy: 0.8553
Epoch 22/40
781/781 [=====] - 49s 63ms/step - loss: 0.5809 - accuracy: 0.8536
Epoch 23/40
781/781 [=====] - 50s 64ms/step - loss: 0.5758 - accuracy: 0.8572
Epoch 24/40
781/781 [=====] - 51s 65ms/step - loss: 0.5781 - accuracy: 0.8570
Epoch 25/40
781/781 [=====] - 51s 65ms/step - loss: 0.5774 - accuracy: 0.8577
Epoch 26/40
781/781 [=====] - 50s 65ms/step - loss: 0.5763 - accuracy: 0.8565
Epoch 27/40
781/781 [=====] - 52s 66ms/step - loss: 0.5779 - accuracy: 0.8545
Epoch 28/40
781/781 [=====] - 51s 65ms/step - loss: 0.5732 - accuracy: 0.8571
Epoch 29/40
781/781 [=====] - 51s 65ms/step - loss: 0.5716 - accuracy: 0.8584
Epoch 30/40
781/781 [=====] - 51s 65ms/step - loss: 0.5717 - accuracy: 0.8595
Epoch 31/40
781/781 [=====] - 52s 66ms/step - loss: 0.5751 - accuracy: 0.8562
Epoch 32/40
781/781 [=====] - 50s 65ms/step - loss: 0.5688 - accuracy: 0.8597
Epoch 33/40
781/781 [=====] - 50s 65ms/step - loss: 0.5725 - accuracy: 0.8584
Epoch 34/40
781/781 [=====] - 50s 65ms/step - loss: 0.5654 - accuracy: 0.8603
Epoch 35/40
781/781 [=====] - 50s 64ms/step - loss: 0.5666 - accuracy: 0.8605
Epoch 36/40
781/781 [=====] - 50s 64ms/step - loss: 0.5637 - accuracy: 0.8594
Epoch 37/40
781/781 [=====] - 50s 64ms/step - loss: 0.5649 - accuracy: 0.8613
Epoch 38/40
781/781 [=====] - 50s 64ms/step - loss: 0.5677 - accuracy: 0.8599
Epoch 39/40
781/781 [=====] - 51s 65ms/step - loss: 0.5699 - accuracy: 0.8599
Epoch 40/40
781/781 [=====] - 50s 64ms/step - loss: 0.5618 - accuracy: 0.8635
```

3.2. Evaluate the model on the test set

Do NOT load the test set until now. Make sure that your model parameters and hyper-parameters are independent of the test set.

```
In [9]: loss_and_acc = model.evaluate(x_test, y_test_vec)

print('loss = ' + str(loss_and_acc[0]))
print('accuracy = ' + str(loss_and_acc[1]))

10000/10000 [=====] - 3s 337us/step
loss = 0.5686637101173401
accuracy = 0.86599996621399
```